



# PostgreSQL for Oracle Users

March 2021



# DUAL

## ORACLE

```
SELECT sysdate FROM dual;
```

## Postgres

```
Select now(); -- transaction  
select clock_timestamp; -- increment
```

```
create table dual as (values(true));
```

# GIVE ME 10 rows

## ORACLE

```
SELECT rownum  
FROM dual  
CONNECT BY rownum <= 10;
```

## Postgres

```
SELECT generate_series  
FROM generate_series(1, 10);
```

Or

```
SELECT ..From a table  
Limit 10;
```

# At least math is easy

# Integer arithmetic

## ORACLE

```
SELECT 1/4 quarter_pounder  
FROM dual;
```

## PostgreSQL

```
SELECT 1/4 royale_with_cheese;
```

# This is madness!

This is *integer arithmetic*!

$10/4 = 2$

$5+10/4 = 7$

$5.0+10/4 = 7.0$

```
SELECT avg(x) FROM foo;  
0
```

$10/4.0 = 2.5$

$5+10/4.0 = 7.5$

$10/4::float = 2.5$

```
SELECT avg(x::float) FROM foo;  
0.22
```

# It's a date

# Date and time types

## ORACLE

`DATE` - Date+time with second resolution

`TIMESTAMP` - Date+time with microsecond resolution

- Optionally can encode time zone

## Postgres

`DATE` - Date with one day resolution

`TIMESTAMP` - Date + time with microsecond resolution



# CREATE a DATE

## ORACLE

```
SELECT to_date('2016/03/26')  
FROM dual;
```

2016/03/26

You **should** specify a date format.

## Postgres

```
SELECT to_date('2016/03/26');
```

Invalid operation: function  
to\_date("unknown") does not  
exist;

You **must** specify a date format.

# Date math

## ORACLE

`date_col1 - date_col2 = number`

SS

## Postgres

`date_col1 - date_col2 = integer`

`tstamp_col1 - tstamp_col2 = interval`

`date_col1 - tstamp_col2 = interval`

`tstamp_col1 - date_col1 = interval`

# Date math

## Oracle

```
date_col + 1 = 2018/04/04  
timestamp_col + 1 = 2018/04/04  
date_col + 0.5 = 2018/04/04 12:00
```

## Postgres

```
date_col + 1 = 2018/04/04  
timestamp_col + '1 day'::interval = 2018/04/04  
date_col + '0.5'::interval = "2018/04/04 12:00:00"
```

# More date fun!

TRUNC(timestamp) returns a DATE...

- ..but only a date by truncating time portion to a date

DATE\_TRUNC() works more like Oracle trunc()

- ...but the parameters are "backward" and
- ...there are a whole new set of datepart abbreviations

`date_trunc('mon', sysdate) = 2016-03-01 00:00:00`

- ...and....

# Weeks start on Monday

## Oracle

```
trunc(t, 'day') = Sunday
```

## Postgres

```
date_trunc('week', d) = Monday
```

```
date_trunc('week', d+1)-1 = Sunday
```

```
date_trunc('week', d+1)-1+6 = Saturday
```

```
select date_trunc('week', now() + '1 day'::interval)-'1  
day'::interval
```

# Looking at strings

# NULL "Handling is Different"

`1 + NULL = NULL`

`sysdate + NULL = NULL`

`SUM(all_null_column) = NULL`

`SUM(some_nulls_column) → Ignore Nulls`

- Oracle

`'string' || NULL = 'string'`

- PostgreSQL

`'string' || NULL = NULL`

# Searching strings

LIKE exists but so does ILIKE

- `lower(foo) LIKE '%interesting%'`
- `foo ILIKE '%interesting%'`

Regex is done by it's own operator

- `lower(foo) ~ '.*interesting.*'`

Use LIKE/ILIKE if you can



# Postgres DDL Standards

## Oracle

```
CREATE TABLE booker.ad_decorations
( advertiser_idNUMBER(38)
, ad_decoration_idNUMBER(38)NOT NULL
, ad_group_idNUMBER(38)NOT NULL
, ad_idNUMBER(38)NOT NULL
, created_byVARCHAR2(8)NOT NULL
, creation_dateDATENOT NULL
, decoration_idNUMBER(38)NOT NULL
, decoration_natural_idVARCHAR2(255)NOT NULL
, decoration_orderNUMBER(10)NOT NULL
, internal_statusNUMBER(10)NOT NULL
, last_updated_byVARCHAR2(8)NOT NULL
, last_updated_dateDATENOT NULL
, marketplace_idNUMBER(38)
, statusNUMBER(10)NOT NULL
, uniqueness_scopeVARCHAR2(64)NOT NULL
, versionNUMBER(38)NOT NULL
)
PCTFREE 10
PCTUSED
INITRANS 8
TABLESPACE bid
STORAGE ( MAXEXTENTS UNLIMITEDPCTINCREASE 0FREELISTS)
/
```

## Postgres

```
CREATE TABLE booker.ad_decorations (
ad_decoration_id NUMERIC(38,0) NOT NULL,
ad_id NUMERIC(38,0) NOT NULL,
decoration_id NUMERIC(38,0) NOT NULL,
marketplace_id NUMERIC(38,0),
advertiser_id NUMERIC(38,0),
decoration_order NUMERIC(10,0) NOT NULL,
status NUMERIC(10,0) NOT NULL,
ad_group_id NUMERIC(38,0) NOT NULL,
decoration_natural_id CHARACTER VARYING(255) NOT
NULL,
uniqueness_scope CHARACTER VARYING(64) NOT NULL,
internal_status NUMERIC(10,0) NOT NULL,
version NUMERIC(38,0) NOT NULL,
created_by CHARACTER VARYING(64) NOT NULL,
creation_date TIMESTAMP(0) WITH TIME ZONE NOT
NULL default (clock_timestamp() at time zone
'UTC'),
last_updated_by CHARACTER VARYING(64) NOT NULL,
last_updated_date TIMESTAMP(0) WITH TIME ZONE
NOT NULL,
program_type NUMERIC(5,0)
)
WITH (
OIDS=FALSE
);
```

```
CREATE INDEX booker.pk_ad_decorations
ON booker.ad_decorations
( ad_decoration_id )
PCTFREE 10
INITTRANS 8
TABLESPACE bid_idx
STORAGE ( MAXEXTENTS UNLIMITEDPCTINCREASE 0 )
ONLINE;
ALTER TABLE booker.ad_decorations
ADD CONSTRAINT pk_ad_decorations
PRIMARY KEY ( ad_decoration_id )
USING INDEX;

CREATE INDEX booker.uq_ads_decors_uq_scope
ON booker.ad_decorations
( ad_group_id, uniqueness_scope, decoration_id )
PCTFREE 10
INITTRANS 8
TABLESPACE bid_idx
STORAGE ( MAXEXTENTS UNLIMITEDPCTINCREASE 0 )
ONLINE;
ALTER TABLE booker.ad_decorations
ADD CONSTRAINT uq_ads_decors_uq_scope
UNIQUE ( decoration_id, uniqueness_scope,
ad_group_id )
USING INDEX;
```

```
ALTER TABLE booker.ad_decorations
ADD CONSTRAINT pk_ad_decorations
PRIMARY KEY (ad_decoration_id);

ALTER TABLE booker.ad_decorations
ADD CONSTRAINT uq_ads_decors_uq_scope
UNIQUE (ad_group_id, uniqueness_scope,
decoration_id );
```

```

CREATE OR REPLACE TRIGGER booker.AUDIT_AD_DECORATIONS
BEFORE INSERT OR UPDATE
ON booker.ad_decorations
FOR EACH ROW
DECLARE
osuserVARCHAR2(8 BYTE);
BEGIN
--
-- get who is making the DML change from the session
--
osuser := session_info.osuser;
--
IF (NOT (dbms_snapshot.i_am_a_refresh)
AND dbms_reutil.replication_is_on) THEN
IF NOT dbms_reutil.from_remote THEN
IF INSERTING THEN
:new.created_by := osuser;
:new.creation_date := sysdate;
:new.last_updated_by := osuser;
:new.last_updated_date := sysdate;
END IF;
--
IF UPDATING THEN
:new.last_updated_by := osuser;
:new.last_updated_date := sysdate;
END IF;
END IF;
END IF;

EXCEPTION

WHEN OTHERS THEN
-- raise error if we cannot set auditing fields
raise_application_error(-20505, 'BOOKER.AUDIT_AD_DECORATIONS
Failed. Rolling Back.' || SQLERRM);
END;
/

```

```

CREATE OR REPLACE FUNCTION booker.audit_trg_function()
RETURNS trigger AS
$BODY$
BEGIN
    IF (TG_OP = 'INSERT') THEN
        NEW.creation_date=clock_timestamp() AT TIME ZONE 'UTC';
        NEW.last_updated_date=clock_timestamp() AT TIME ZONE 'UTC';
        NEW.created_by=SESSION_USER::TEXT;
        NEW.last_updated_by=SESSION_USER::TEXT;
        RETURN NEW;
    ELSIF (TG_OP = 'UPDATE') THEN
        NEW.last_updated_date=clock_timestamp() AT TIME ZONE 'UTC';
        NEW.last_updated_by=SESSION_USER::TEXT;
        RETURN NEW;
    ELSE
        RAISE WARNING 'audit_trg_function exeception in %.% - OTHER ACTION
OCCURRED: %, AT %',TG_TABLE_SCHEMA,TG_TABLE_NAME,TG_OP,clock_timestamp();
        RETURN NULL;
    END IF;
EXCEPTION
    WHEN DATA_EXCEPTION THEN
        RAISE WARNING 'audit_trg_function exeception in %.% - UDF ERROR [DATA
EXCEPTION\] - SQLSTATE: %, SQLERRM:
%',TG_TABLE_SCHEMA,TG_TABLE_NAME,SQLSTATE,SQLERRM;
        RETURN NULL;
    WHEN UNIQUE_VIOLATION THEN
        RAISE WARNING 'audit_trg_function exeception in %.% - UDF ERROR [UNIQUE\]
- SQLSTATE: %, SQLERRM: %',TG_TABLE_SCHEMA,TG_TABLE_NAME,SQLSTATE,SQLERRM;
        RETURN NULL;
    WHEN OTHERS THEN
        RAISE WARNING 'audit_trg_function exeception in %.% - UDF ERROR [OTHER\]
- SQLSTATE: %, SQLERRM: %',TG_TABLE_SCHEMA,TG_TABLE_NAME,SQLSTATE,SQLERRM;
        RETURN NULL;
END;
$BODY$
LANGUAGE plpgsql VOLATILE
COST 100;
CREATE TRIGGER AUDIT_TRIGGER BEFORE INSERT OR UPDATE ON booker.ad_decorations
FOR EACH ROW EXECUTE PROCEDURE booker.audit_trg_function();

```

# Postgres unexpected behavior

Explicit rollback required on SQLException (when autocommit is off)

- Use savepoints to allow rolling back individual items in the same commit
- Oracle/Mysql have an implicit rollback on SQLException

TIMESTAMP handling

- TIMESTAMP (without time zone) stores the time in the time zone of the client
- Use TIMESTAMP WITH TIME ZONE and don't provide a time zone to force storing timestamp in UTC

# PostgreSQL vs Oracle Play Book

	ORACLE	Postgres
<a href="#">link</a>	Common Data Types	Common Data Types
<a href="#">link</a>	DB hints	Query Planning
<a href="#">link</a>	Inline Views	Inline Views
<a href="#">link</a>	Merge	SQL Merge
<a href="#">link</a>	DBMS Random	Random()
<a href="#">link</a>	DBMS Output	Raise
<a href="#">link</a>	Procedure and Function	Functions
<a href="#">link</a>	OLAP Function	Windowing Function
<a href="#">link</a>	Execute Immediate	Execute
<a href="#">link</a>	Index Organized Table	Postgres Cluster Tables
<a href="#">link</a>	Table Constraints	Table Constraints
<a href="#">link</a>	Table Partitioning including: RANGE, LIST, HASH, COMPOSITE, Automatic LIST	Table Partitioning including: RANGE, LIST
<a href="#">link</a>	Temporary Tables	Temporary Tables
<a href="#">link</a>	Unused Columns	Drop Columns

	Oracle	Postgres
<a href="#">link</a>	Read Only Tables	Read only Roles
<a href="#">link</a>	Index Types	Index Types
<a href="#">link</a>	B Tree Index	B Tree Index
<a href="#">link</a>	Composite Index	Composite Index
<a href="#">link</a>	Bitmapped Index	BRIN Index
<a href="#">link</a>	Functional Index	Expression Index
<a href="#">link</a>	Global and Local Index	Local Index
<a href="#">link</a>	LOB and Secure Files	LOB
<a href="#">link</a>	Materialized Views	Materialized Views
<a href="#">link</a>	Oracle Triggers	Postgres Functional Triggers
<a href="#">link</a>	Views	Views
<a href="#">link</a>	Sequences	Sequences
<a href="#">link</a>	Database Links	Postgres DBlink and FDWrapper



# Questions ?

# Thank you !